NumeroLogic: Number Encoding for Enhanced LLMs' Numerical Reasoning

Eli Schwartz¹, Leshem Choshen^{2,3}, Joseph Shtok¹, Sivan Doveh¹, Leonid Karlinsky², Assaf Arbelle¹

¹IBM Research, ²MIT-IBM Watson AI Lab, ³MIT

Abstract

Language models struggle with handling numerical data and performing arithmetic operations. We hypothesize that this limitation can be partially attributed to non-intuitive textual numbers representation. When a digit is read or generated by a causal language model it does not know its place value (e.g. thousands vs. hundreds) until the entire number is processed. To address this issue, we propose a simple adjustment to how numbers are represented by including the count of digits before each number. For instance, instead of "42", we suggest using "2:42" as the new format. This approach, which we term NumeroLogic, offers an added advantage in number generation by serving as a Chain of Thought (CoT). By requiring the model to consider the number of digits first, it enhances the reasoning process before generating the actual number. We use arithmetic tasks to demonstrate the effectiveness of the NumeroLogic formatting. We further demonstrate NumeroLogic applicability to general natural language modeling. improving language understanding performance in the MMLU benchmark.

1 Introduction

Large Language Models (LLMs) struggle with numerical and arithmetical tasks. Despite continuous improvements, even the most advanced models like GPT-4 [1] still exhibit poor performance when con-



Figure 1: Reading numbers in a causal manner from left to right is sub-optimal for LLMs, as it is for humans. The model has to reach the final digits of the number before it can infer the place value of the first digit. To address this issue we propose "NumeroLogic", a numerical format where the digit count is indicated before the actual number. Image adapted from imagery generated by DALL-E 3 [2].

fronted with tasks such as multiplying 3-digit numbers [13]. Recent studies ([10, 13]) have proposed techniques to improve arithmetic in LLMs, such as the Chain of Thought (CoT; [16]) method, which pushes the model to anticipate the entire sequence of algorithmic steps rather than just the final output. While these strategies offer valuable insights into the capabilities of LLMs, they primarily concentrate on post-hoc solutions for specific arithmetic challenges and do not present a practical solution for pretraining LLMs. Our research, however, focuses on solutions applicable to self-supervised language modeling in general, utilizing arithmetic exercises primarily for evaluating their impact.

We hypothesize that one of the challenges LLMs face when dealing with numerical tasks is the textual representation of numbers. In today's most popular decoder-based LLMs, each token attends only to previous tokens. When a model "reads" a token representing a digit (or multiple digits) it cannot tell its place value, i.e. '1' can represent 1 million, 1 thousand, or a single unit. Only when reaching the end of the number might the model update its representation of the previous digit tokens to be related to their real place value.

To address this issue, we propose a straightforward reformatting technique called "NumeroLogic," which involves adding the number of digits as a prefix to numbers. This lets the model know in advance what is the place value of a digit before it is read. This simple change also offers another benefit, when the model is generating a number it needs to first reason about what is going to be the number of digits. This acts as a Chain of Thought (CoT) [16], encouraging the model to perform some reasoning before it begins to predict digits. Implementing the suggested reformatting does not necessitate any alterations to the model's architecture; it can be accomplished through text pre- and post-processing based on regular expressions.

We demonstrate that NumeroLogic enhances the numerical abilities of LLMs across both small and larger models (up to 7B parameters). This enhancement is showcased through supervised training on arithmetic tasks and its application in self-supervised causal language modeling to enhance general language comprehension.

2 Related Work

Recently, there has been a significant interest in enhancing the numerical capabilities of LLMs. One approach to investigating these capabilities is by assessing their performance in arithmetic tasks. Several recent studies have proposed methods to enhance performance in these tasks. One strategy involves reversing the expected result order from the least significant digit to the most significant [10]. Another strategy is using an elaborated CoT where the model is taught to predict all steps of an algorithm predefined for each arithmetic task [10]. In [13], it is noted that the model learns to rely too heavily on positional encoding when trained for a specific arithmetic task. They suggest ways to overcome it, e.g. adding random whitespaces in the middle of numbers. These studies aim to enhance the performance of arithmetic tasks by offering tailored solutions to the associated challenges. In contrast, our focus is on identifying solutions that benefit general language modeling rather than just arithmetic tasks, with arithmetic tasks being used solely for measuring improvements.

Another aspect important for LLMs numerical capabilities is the tokenization process. The commonly used Byte Pair Encoding (BPE) based methods [4, 12] for tokenization are based on the corpus distribution and can split a number to tokens in unintuitive ways. Different foundation models took different approaches when dealing with number tokenization. PaLM [3], Llama [15], and Mistral [8] force each digit to have a single token. GPT-3.5 and GPT-4 define a token for each up to 3-digit number [1]. Somewhat related to our work, in [14], they highlighted an issue with the GPT approach. They show that dividing large numbers into 3-digit segments from left to right undermines arithmetic performance. They suggest overcoming it by artificially inserting commas between digits to control the splitting.

3 NumeroLogic

We introduce NumeroLogic, a technique for boosting causal LLM's numerical capabilities. The concept involves adding a digit count before numbers, enabling

the model to know the place values of digits before reaching the final digits of a number. Additionally, the model needs to predict the total number of digits before generating a number, acting as a simplified CoT, prompting it to reason about the number that is going to be generated.

We add special tokens to help represent numbers with the number-of-digit prefix, "<startnumber>", "<midnumber>", and "<endnumber>" (or, for simplicity, "<sn>", "<mn>", and "<en>"). For floating points, the prefix includes both the number of digits of the integer part and the decimal part. For example, "42" is replaced by "<sn>2<mn>42<en>" and "3.14" is replaced by "<sn>1.2<mn>3.14<en>". When using the LLM to generate numbers, we disregard the information about the number of digits and only retain the generated number itself. Although not within the scope of this study, it may be feasible to leverage the additional information to identify discrepancies, wherein the model predicts a certain digit count but produces a number with a different count of digits.

The NumeroLogic approach includes basic text preprocessing and post-processing steps that occur before and after the tokenizer's encoding and decoding methods, respectively. Both can be implemented based on regular expressions:

```
def preprocess_all_numbers(text):
  def f (match):
    num = match.group(0)
    i = match.group(1)
    li = len(i)
    d = match.group(3)
    ld = len(d) if d else 0
    if d:
       prefix = f' < sn > \{li\}.\{ld\} < mn > \}
    else:
       prefix = f' < sn > \{li\} < mn > \}
    return prefix + num + '<en>'
  pattern = ((d+)(((d+))?)
  return re.sub(pattern, f, text)
def postprocess_all_numbers(text):
  pattern = \langle sn \rangle [\d\.] + \langle mn \rangle
  text = re.sub(pattern, '', text)
  text = re.sub('<en>', '', text)
  return text
```

For small transformers (NanoGPT [9]), we train all parameters from scratch with character-level tokenization. For small transformers, we also replace the special tokens with single characters, "<sn>", "<mn>", and "<en>" are replaced with "{", ":", and "}", respectively. For larger transformers, we start from pre-trained models. We add the new special tokens to the tokenizer's vocabulary and expand the embedding layer and the final fully connected layer to fit the new vocabulary size. When continuing training on causal language modeling or fine-tuning on supervised arithmetic tasks, we use low-rank adaptation (LoRA) [7]. We apply LoRA for the attention block projection matrices (Q, K, V, O) and train the modified embedding layer and the final fully-connected layer in full rank.

4 Experiments

To test the effect of NumeroLogic we conducted several experiments. First, we tested supervised training of a small language model (NanoGPT) on various arithmetic tasks. We then test the scalability to larger models (Llama2-7B). Finally, we test self-supervised pretraining of Llama2-7B, with the suggested formatting, and test on general language understanding tasks.

4.1 Arithmetic tasks with small model

We trained NanoGPT [9] from scratch in a supervised manner jointly on 5 arithmetic tasks: addition, subtraction, multiplication, sine, and square root. Addition and subtraction are performed with up to 3-digit integer operands. Multiplications are performed with up to 2-digit integer operands. Sine and square root with 4 decimal-places floating point operands and results. The operand range for sine is within $[-\pi/2, \pi/2]$. The operand range for the square root is within [0, 10]. The model is trained in a multi-task fashion on all 5 tasks, with 10K training samples for each task except for multiplication, for which 3K samples are used. We followed the protocol from Section D.2 in [10].

Figure 1 compares the results of training with plain

Op.	Num. digit	int/ float	Plain	Numero Logic	Gain
+	3	int	88.37	99.96	+11.6
_	3	$_{ m int}$	73.76	97.20	+23.4
*	2	int	13.81	28.94	+15.1
sine	4	float	30.59	34.59	+4.00
sqrt	4	float	22.13	26.66	+4.53

Table 1: NanoGPT arithmetic tasks accuracy with NumeroLogic encoding. A single model is jointly trained for all tasks. The encoding produces high accuracy gains for all tasks.

numbers to training with the NumeroLogic encoding. For addition and subtraction, a model trained with plain numbers reached 88.37% and 73.76% accuracy, respectively, while with the NumeroLogic encoding, the tasks are almost solved (99.96% and 97.2%). For multiplication, we observe more than doubling of the accuracy, from 13.81% to 28.94%. Furthermore, for the floating point operations, sine and square root, we see a significant improvement of 4% for both tasks.

4.2 Arithmetic tasks with larger model

Next, we test how the method scales to a larger model. For this experiment, we fine-tune a pretrained Llama2-7B model [15]. In this experiment, we again tested the same five arithmetic tasks: addition, subtraction, multiplication, sine, and square root. For addition (5 digit), subtraction (5 digit), and multiplication (3 digit) we tested on two versions - integers and floating point numbers. For generating a random Ndigit floating point operand we first sample an up to N-digit integer and then divide it by a denominator uniformly sampled from $\{10^0, 10^1, ..., 10^N\}$. For each of the addition, subtraction, and multiplication tasks, we generated 300K random equations as a training set. The sine and square root operands and results are generated with 5 decimal place accuracy, we generated 30K random equations for the training sets of these tasks. Since we are working with a pretrained model we add new tokens ("<sn>", "<mn>", and "<en>") to the tokenizer' vocabulary. We finetune one model per

	Num.	int/		Numero	
Op.	digit	float	Plain	Logic	Gain
+	5	int	99.86	100.0	+0.14
_	5	int	99.60	99.93	+0.33
*	3	int	34.20	35.33	+1.13
+	5	float	91.40	94.43	+3.03
_	5	float	88.76	92.73	+3.97
*	3	float	24.73	31.03	+6.30
$_{ m sine}$	5	float	25.06	28.13	+3.07
sqrt	5	float	13.00	17.16	+4.16

Table 2: Llama2-7B arithmetic tasks accuracy with NumeroLogic encoding. We observe significant gains thanks to the NuemroLogic encoding for all tasks where performance is not saturated.

task with LoRA [7] (rank 8), we also train in full-rank the embedding layer and the final fully-connected layer since their parameters are extended to accommodate the larger vocabulary size.

The results are presented in Table 2. Addition and subtraction of integers are mostly solved by a model as large as Llama2-7B even for much larger numbers (e.g. 20-digit). For our 5-digit experiments, the plain text baselines reached 99.86% and 99.6% performance, for addition and subtraction, respectively. Despite the high performance of plain text, we still observe an improvement when using NumerLogic, with a perfect 100% for addition and rectification of more than 80% of the subtraction mistakes, reaching 99.93% accuracy for subtraction. For all other, non-saturated, tasks we observed significant gains of 1%-6%.

4.3 Self-Supervised Pretraining

Our approach differs from other methods in that it is not specialized for a specific task, such as arithmetic, but rather designed for general language modeling tasks involving text with numerical values. To test this capability we continue the pretraining of LLama2-7B with the causal text modeling objective (next token prediction). We train on text from the RefinedWeb dataset [11]. The goal is to teach the model to read and write numbers in the NumeroLogic format with-

	Change
Social sciences	+0.1%
Humanities	+0.43%
STEM	+0.79%
Others	+1.19%

Table 3: MMLU accuracy change due to NumeroLogic encoding on tasks from different fields. STEM tasks which are more likely to require numerical understanding enjoy higher improvement.

	Change
Tasks with numbers Tasks without numbers	$+1.16\% \\ +0.14\%$

Table 4: MMLU accuracy change due to NumeroLogic encoding on tasks with and without numbers. Tasks with numbers enjoy higher improvement.

out forgetting its previously acquired knowledge. To facilitate this, we perform the continued pretraining with LoRA. We then test the model in a zero-shot manner on the massive multitask language understanding benchmark (MMLU) [6, 5].

In Figure 2, we present the MMLU 0-shot results obtained from training the model using plain numbers versus NumeroLogic encoding on an equal number of tokens. While training with plain numbers does not enhance the model's accuracy compared to the pretrained model, employing NumeroLogic encoding results in a statistically significant improvement of 0.5%. The MMLU benchmark encompasses tasks from diverse domains, some emphasizing analytical skills and numerical comprehension while others do not. In Table 3, we delve into the impact of NumeroLogic on MMLU tasks categorized by field. As anticipated, tasks in STEM fields exhibit more substantial enhancements compared to those in social sciences and humanities. Table 4 provides a detailed analysis of NumeroLogic's performance boost across tasks containing numbers versus those that do not. Consistently, tasks involving numbers show a more pronounced improvement.

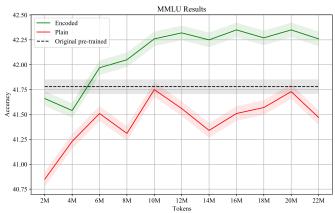


Figure 2: MMLU Accuracy of Llama2-7B. Continuing self-supervised pretraining on web-curated text tokens, when numbers are encoded with NumeroLogic, helps improve the performance beyond the pretrained model or a model trained on the same text with plain numbers.

		Operands		
ند		Plain	Encoded	
Sesult	Plain	88.37%	98.05%	
m Re	Encoded	89.34%	99.78%	

Table 5: Testing the effect of encoding the equation's operands vs. result. Tested on the addition task with NanoGPT. Either encoding the operands (i.e. input comprehension) or encoding the results (i.e. CoT effect) have a positive effect, with a stronger effect for operands' encoding. Encoding both the operands and the result provides the best performance.

Encoding	3-digit Multiplication
Plain (e.g. "100")	34.20%
Multi special tokens ("<3digitnumber>100")	33.56%
Only prefix (" <sn>3<mn>100")</mn></sn>	34.93%
NumeroLogic (" <sn>3<mn>100<en>")</en></mn></sn>	35.33%

Table 6: Different encoding alternatives

4.4 Ablation studies

4.4.1 Encoding operands vs. results

We experimented to test the effect of operand encoding vs. the expected output (equation result) encoding. Operand encoding primarily influences the model's comprehension of numerical values in the input, while result encoding is more associated with CoT, prompting the model first reason about the expected number of digits. We repeat the experiment from Section 4.1, but with the NumeroLogic encoding applied only to the operands or to the results and report the 3-digit addition results for the different variants. The results are presented in Table 5. We find that both operands and results encodings are beneficial, with a stronger impact attributed to encoding the results. Applying NumeroLogic to all numbers, both operands and results, yields the highest level of accuracy.

4.4.2 Different Encodings

We experimented with different formats for providing the number of digits. One alternative we tested is defining a set of new special tokens representing each possible number of digits, {<1digitnumber>, <2digitnumber>,...}. We observed that the performance of having multiple special tokens is even lower than plain numbers. It might be due to the unbalanced distribution of numbers. E.g. numbers with a single digit are much less frequent in a data of 3-digit additions, it is possible the model has not seen enough single-digit numbers to learn a good representation of the <1digitnumber> token. Another alternative we tested is removing the "end of number" token (<en>), keeping only the number prefix, e.g. "<sn>3<mn>100". This works better than plain but slightly worse than the full NumeroLogic encoding. The results are summarized in Table 6.

5 Conclusions

In this paper, we introduced NumeroLogic, an innovative approach to enhancing language models' comprehension and generation of numerical data. Our hypothesis centered around the challenge of numerical representation in text for large language models (LLMs), which traditionally struggle with arithmetic and numerical understanding. By presenting numbers with a prefixed notation indicating their digit count, we proposed a solution that aids models in recognizing the place value of digits before the complete presentation of the full number. This method not only facilitates a better understanding of numbers but also prompts the model to reason about the magnitude of numbers it intends to generate, effectively integrating a form of Chain of Thought (CoT) reasoning.

Our experiments spanned from arithmetic tasks to general natural language understanding, using both small and large model architectures. The results consistently demonstrated that NumeroLogic significantly improves models' numerical capabilities. In arithmetic tasks, we observed marked improvements in accuracy across different tasks including integer and floating-point operations. This improvement was not just confined to specialized arithmetic tasks but also extended to general language modeling, as evidenced by performance boosts in the MMLU benchmark for tasks requiring numerical understanding.

In conclusion, NumeroLogic offers a simple yet powerful tool for enhancing the numerical capabilities of language models without necessitating architectural modifications. Its success across various tasks and models highlights its potential as a general-purpose solution, particularly beneficial in domains where numerical understanding is crucial.

References

- [1] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. arXiv preprint arXiv:2303.08774, 2023.
- [2] James Betker, Gabriel Goh, Li Jing, Tim Brooks, Jianfeng Wang, Linjie Li, Long Ouyang, Juntang Zhuang, Joyce Lee, Yufei Guo, et al. Improving image generation with better captions. Computer Science. https://cdn. openai. com/papers/dall-e-3. pdf, 2(3):8, 2023.
- [3] Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, et al. Palm: Scaling language modeling with pathways. *Journal of Machine Learning Research*, 24(240):1–113, 2023.
- [4] Philip Gage. A new algorithm for data compression. *The C Users Journal*, 12(2):23–38, 1994.
- [5] Dan Hendrycks, Collin Burns, Steven Basart, Andrew Critch, Jerry Li, Dawn Song, and Jacob Steinhardt. Aligning ai with shared human values. Proceedings of the International Conference on Learning Representations (ICLR), 2021.
- [6] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. Proceedings of the International Conference on Learning Representations (ICLR), 2021.
- [7] Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. arXiv preprint arXiv:2106.09685, 2021.
- [8] Albert Q Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna

- Lengyel, Guillaume Lample, Lucile Saulnier, et al. Mistral 7b. arXiv preprint arXiv:2310.06825, 2023.
- [9] Andrej Karpathy. Nanogpt. https://github.com/karpathy/nanoGPT, 2022.
- [10] Nayoung Lee, Kartik Sreenivasan, Jason D. Lee, Kangwook Lee, and Dimitris Papailiopoulos. Teaching arithmetic to small transformers. In The Twelfth International Conference on Learning Representations, 2024.
- [11] Guilherme Penedo, Quentin Malartic, Daniel Hesslow, Ruxandra Cojocaru, Alessandro Cappelli, Hamza Alobeidli, Baptiste Pannier, Ebtesam Almazrouei, and Julien Launay. The RefinedWeb dataset for Falcon LLM: outperforming curated corpora with web data, and web data only. arXiv preprint arXiv:2306.01116, 2023.
- [12] Rico Sennrich, Barry Haddow, and Alexandra Birch. Neural machine translation of rare words with subword units. arXiv preprint arXiv:1508.07909, 2015.
- [13] Ruoqi Shen, Sébastien Bubeck, Ronen Eldan, Yin Tat Lee, Yuanzhi Li, and Yi Zhang. Positional description matters for transformers arithmetic. arXiv preprint arXiv:2311.14737, 2023.
- [14] Aaditya K Singh and DJ Strouse. Tokenization counts: the impact of tokenization on arithmetic in frontier llms. arXiv preprint arXiv:2402.14903, 2024.
- [15] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. Llama: Open and efficient foundation language models. arXiv preprint arXiv:2302.13971, 2023.
- [16] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. Advances in neural information processing systems, 35:24824–24837, 2022.